

Zap the Flakes!

Leveraging AI to Combat Flaky Tests with CANNIER

devtalks.ro '25

[Daniel Hiller](#)



Red Hat



agenda

- about me
- about flakes
- pre-merge-detection v1
- CANNIER
- pre-merge-detection v2
- Q&A



about me

- Software Engineer | [Red Hat OpenShift Virtualization](#)
- [KubeVirt](#) | CI & automation in general
- [prow.ci.kubevirt.io](#)
- [CI-Health](#)
- Flaky Tests



[kubevirt.io](#) - Virtualization for Kubernetes



Dev<Talks/>

Red Hat OpenShift Virtualization



about flakes

PR History: kubevirt/kubevirt #9445

	9d41878		
pull-kubevirt-e2e-k8s-1.25-sig-compute-migrations	1637934812398358528	1636633531918585856	1636403749595385856
pull-kubevirt-e2e-k8s-1.25-sig-compute	1637934815221125120	1636403757321293824	
pull-kubevirt-e2e-k8s-1.25-sig-network	1637934813975416832	1636403756985749504	
pull-kubevirt-e2e-k8s-1.25-sig-operator	1637934815393091584	1636403757992382464	
pull-kubevirt-e2e-k8s-1.25-sig-storage	1637934814088663040	1636633532048609280	1636403756704731136
pull-kubevirt-e2e-k8s-1.26-sig-compute	1637934816471027712	1636404222087925760	
pull-kubevirt-e2e-k8s-1.26-sig-network	16379348160853151744	1636403758015120844	

pull-kubevirt-e2e-k8s-1.25-sig-compute-migrations #1636633531918585856

Job History PR History Artifacts

Test started last Friday at 8:40 AM **passed** after 1h9m49s. ([more info](#))

JUnit

91/1406 Tests Passed!

1315/1406 Tests Skipped.

Build Log

Show all hidden lines [Raw build-log.txt](#)

pull-kubevirt-e2e-k8s-1.25-sig-compute-migrations #1636403749595385856

Job History PR History Artifacts

Test started last Thursday at 7:33 PM **failed** after 1h18m18s. ([more info](#))

JUnit

1/1406 Tests Failed.

Tests Suite: [rfe_id:393][crit:high][vendor:cnv-qe@redhat.com][level:system][sig-compute] VM Live Migration [Serial]with a dedicated migration network Should migrate over that network 4m18s

90/1406 Tests Passed!

1315/1406 Tests Skipped.

source: <https://prow.ci.kubevirt.io/pr-history/?org=kubevirt&repo=kubevirt&pr=9445>



about flakes

a *flake*

is a **test** that

without any code change

will either **fail** or **pass** in successive runs



about flakes

“... test flakiness was a **frequently encountered problem**, with

- **20%** of respondents claiming to experience it **monthly**,
- **24%** encountering it on a **weekly** basis and
- **15%** dealing with it **daily**”

source: [“A survey of flaky tests”](#)



about flakes

“... In terms of severity, of the **91% of developers** who claimed to deal with flaky tests at least a few times a year,

- **56%** described them as a **moderate** problem and
- **23%** thought that they were a **serious** problem. ...”

source: [“A survey of flaky tests”](#)



impact of flakes

from [“A survey of flaky tests”](#):

- **97%** of flakes were **false alarms***, and
- **more than 50%** of flakes could not be **reproduced** in isolation

conclusion: “ignoring flaky tests is ok”

*code under test actually is not broken, but it works as expected



impact of flakes



impact of flakes

in CI automated testing **MUST** give a **reliable signal of stability**

any failed test run signals that the product is **unstable**

test runs failed due to flakes do not give this reliable signal

they only waste time



impact of flakes

Flaky tests waste everyone's time!

For the **project community** they cause

- **slowdown of merging** pull requests - “retest trap”
- **reversal of acceleration** effects (i.e. batch testing)
- **waste** of CI resources



impact of flakes

Flaky tests waste everyone's time!

For **individual contributors** they cause

- **prolonged feedback cycles**
- **test trust issues**



pre-merge-detection

Goal: flaky tests **must not** enter the code base!



pre-merge-detection v1

Initial approach:

- look at files touched in pull request
- determine changed files containing e2e tests
- run all contained tests **5 times** in **random** order



pre-merge-detection v1

[check-tests-for-flakes test lane](#)

why: catch flakes before entering
main

([source](#))

```
ginkgo_params="$ginkgo_params -no-color -succinct --label-filter=!QUARANTINE -randomize-all"
for test_file in $(echo "${NEW_TESTS}" | tr '|' '\n'); do
    ginkgo_params+=" -focus-file=${test_file}"
done
|
echo "Test lane: ${TEST_LANE}, preparing cluster up"

if [[ ! "$ginkgo_params" =~ -dry-run ]]; then
    make cluster-up
    make cluster-sync
else
    # Ginkgo only performs -dryRun in serial mode.
    export KUBEVIRT_E2E_PARALLEL="false"
    NUM_TESTS=1
fi

for i in $(seq 1 "$NUM_TESTS"); do
    echo "Test lane: ${TEST_LANE}, run: $i"
    if ! FUNC_TEST_ARGS="$ginkgo_params" make functest; then
        echo "Test lane: ${TEST_LANE}, run: $i, tests failed!"
        exit 1
    fi
done
```



pre-merge-detection v1

why:


- eliminate order-dependencies
- running five times is the boundary to find 88% of flaky tests



pre-merge-detection v1

initial situation:

- most e2e tests take 10sec - 2mins to run
- 5 times re-run has “only” 88% chance of detection



pre-merge-detection v1

problems discovered with v1 approach:

- shotgun approach causes too many tests to be run
 - capping amount of tests re-run required to avoid timeouts
 - CI users complain about failing tests they didn't touch
- re-run lane takes ~1h on average



CANNIER

CANNIER is an approach
for **reducing the time cost** of rerunning-based detection techniques
by **combining them with machine learning models**

source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



CANNIER


“... we found that CANNIER was able to reduce the time cost (and therefore monetary cost) *[of re-running tests]* by an **average of 88% ...**”

source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



CANNIER

single model



source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



CANNIER

feature set (complete)

#	Feature	Description
1	Read Count	Number of times the filesystem had to perform input [9].
2	Write Count	Number of times the filesystem had to perform output [9].
3	Run Time	Elapsed wall-clock time of the whole test case execution.
4	Wait Time	Elapsed wall-clock time spent waiting for input/output operations to complete.
5	Context Switches	Number of voluntary context switches.
6	Covered Lines	Number of lines covered.
7	Source Covered Lines	Number of lines covered that are not part of test cases.
8	Covered Changes	Total number of times each covered line has been modified in the last 75 commits.
9	Max. Threads	Peak number of concurrently running threads.
10	Max. Children	Peak number of concurrently running child processes.
11	Max. Memory	Peak memory usage.
12	AST Depth	Maximum depth of nested program statements in the test case code.
13	Assertions	Number of assertion statements in the test case code.
14	External Modules	Number of non-standard modules (i.e., libraries) used by the test case.
15	Halstead Volume	A measure of the size of an algorithm's implementation [21, 57, 59].
16	Cyclomatic Complexity	Number of branches in the test case code [39, 57, 59].
17	Test Lines of Code	Number of lines in the test case code [57, 59].
18	Maintainability	A measure of how easy the test case code is to support and modify [19, 71].

source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



CANNIER

feature set (runtime)

1 Read Count	Number of times the filesystem had to perform input [9].
2 Write Count	Number of times the filesystem had to perform output [9].
3 Run Time	Elapsed wall-clock time of the whole test case execution.
4 Wait Time	Elapsed wall-clock time spent waiting for input/output operations to complete.
5 Context Switches	Number of voluntary context switches.
6 Covered Lines	Number of lines covered.
7 Source Covered Lines	Number of lines covered that are not part of test cases.
8 Covered Changes	Total number of times each covered line has been modified in the last 75 commits.
9 Max. Threads	Peak number of concurrently running threads.
10 Max. Children	Peak number of concurrently running child processes.
11 Max. Memory	Peak memory usage.

source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



CANNIER

feature set (static analysis)

12	AST Depth	Maximum depth of nested program statements in the test case code.
13	Assertions	Number of assertion statements in the test case code.
14	External Modules	Number of non-standard modules (i.e., libraries) used by the test case.
15	Halstead Volume	A measure of the size of an algorithm's implementation [21, 57, 59].
16	Cyclomatic Complexity	Number of branches in the test case code [39, 57, 59].
17	Test Lines of Code	Number of lines in the test case code [57, 59].
18	Maintainability	A measure of how easy the test case code is to support and modify [19, 71].

source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



pre-merge-detection v2.0

questions

- CANNIER
 - implemented in Python
 - KubeVirt uses Go
- Runtime Data
 - where to store
 - when to capture
- data science
 - Python has well known frameworks
 - Go state unsure




pre-merge-detection v2.0

algorithm

steps:

- gather the set of changed tests
- per each element
 - *run once to get the runtime data*
 - extract the static analysis data
 - fetch the prediction using entire feature vector
 - add to the set of re-run tests if in iDFClass
- re-run the reduced set of tests



source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>




pre-merge-detection v2.0

implementation

parts:

- code (Go)
 - test set extraction
 - feature extraction
 - model prediction
 - model generation
- test lane (Bash)
- model deployment (YAML)



source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



pre-merge-detection v2.1

improvements

parts:

- model
 - generation
 - deployment
- test lane
- code
 - test set extraction
 - feature extraction
 - model prediction



parts:

- model
 - ...
 - **automatic updates**
- ~~test lane~~ -> **prowl external-plugin**
 - runs on presubmits
 - runs on postcommits (probably with a larger test set)
 - adds helpful feedback
- code
 - test set extraction
 - feature extraction
 - model prediction



pre-merge-detection v2.1

prowl external-plugin

- runs on presubmits
- adds helpful feedback
- gives advice according to the feature set what can be improved



pre-merge-detection v2.2

improvements

- components of feature vector provide insightful advice to the contributor
 - i.e. high cyclomatic complexity advises to reduction etc.
 - therefore it's valuable to attach the analysis data to the PR
- possibly increase number of reruns (>5)
 - reduced runtime overall leaves time for more reruns
 - feature vector contains runtimes, thus we can estimate the total re-run time better and optimize for it, i.e. group tests by runtime classes

source: <https://www.gregorykapfhammer.com/research/papers/parry2023/>



pre-merge-detection v1.5 (the present)

approach:

- **gather the exact set of changed tests**
- run these 5 times in random order



Links

- KubeVirt resources

- Initial pull request (draft): <https://github.com/kubevirt/project-infra/pull/3930>
- Presentation Squash The Flakes @ FOSDEM '24:
<https://archive.fosdem.org/2024/schedule/event/fosdem-2024-1805-squash-the-flakes-how-to-minimize-the-impact-of-flaky-tests/>

- CANNIER

- paper: <https://www.gregorykapfhammer.com/research/papers/parry2023/>
- implementation: <https://github.com/flake-it/cannier-framework/>



Q&A

Any questions?

Any suggestions for improvement?

Who else is trying to tackle this problem?

What have you done to solve this?

download slides:



[slides.pdf](#)



Thank you for attending!

Further questions?

Feel free to send questions and comments to:

dhiller@redhat.com

 [kubernetes.slack.com/
@dhiller](https://kubernetes.slack.com/@dhiller)

 [@dhiller@fosstodon.org](https://fosstodon.org/@dhiller)

dhiller.dev

contact me:



kubevirt.io

Virtualization for Kubernetes

KubeVirt welcomes all kinds of contributions!

- **Weekly community meeting every Wed 3PM CET**
- Links:
 - [KubeVirt website](https://kubevirt.io)
 - [KubeVirt user guide](#)
 - [KubeVirt Contribution Guide](#)
 - [GitHub](#)
 - Kubernetes Slack channels
 - [#virtualization](#)
 - [#kubevirt-dev](#)

Dev<Talks/>

